

# CS 156a FINAL

Cason Shepard

January 18, 2023

## problem 1

We will use the formula  $d = \frac{Q(Q+3)}{2}$ . Since  $Q = 10$ , we get  $d = \frac{130}{2} = 65$ . This is not an option.

The answer is [e]

## problem 2

Consider the average of two logistic functions. Sometimes the results will simply not be logistic.

The answer is [d]

## problem 3

In this instance, d is false. This is because we cannot generalize to all instances. When over fitting occurs, we have a hypothesis that generates a very small  $E_{in}$ , but a large  $E_{out}$ . Consider we have a small training set, but a very large test set. In this case we can expect our hypothesis function to be very well fitted to the training data, thus  $E_{in}$  will be very small. However, due to our small training set, the  $E_{out}$  could be very large. In this case, we might be wary that our data is over fitted, but a small training set prevents this determination. Thus, we cannot *always* determine if the data is over fitted through comparing  $E_{in}$  and  $E_{out}$ .

The answer is [d]

## problem 4

Stochastic noise is random noise that is a part of the target function we are trying to evaluate, where as we view deterministic noise as a byproduct of the hypothesis we choose. Thus, stochastic noise is independent of the hypothesis set.

The answer is [d]

## problem 5

We know that  $w_{reg}^T w_{reg} = C$ . So, from the textbook, we know that  $w_{reg} = w_{lin}$  to satisfy the equation.

The answer is [a]

## problem 6

It is known that augmented error constraints can be used to measure the goodness of a soft-order constraint. Thus, soft-order constraints can be translated into augmented error.

The answer is [b]

## problem 7

After running the code attached, the lowest  $E_{in}$  was found for the 8-versus-all classifier.

(a)  $E_{in} = 0.076$

(b)  $E_{in} = 0.091$

(c)  $E_{in} = 0.088$

(d)  $E_{in} = 0.074$

(e)  $E_{in} = 0.088$

The answer is [d]

problem 8

After running the code attached, the lowest  $E_{out}$  was found for the 1-versus-all classifier.

- (a)  $E_{out} = 0.107$
- (b)  $E_{out} = 0.022$
- (c)  $E_{out} = 0.099$
- (d)  $E_{out} = 0.083$
- (e)  $E_{out} = 0.100$

The answer is [b]

problem 9

Using the 5-versus-all classifier, the following  $E_{out}$  was produced:

Non-Transformed:  $E_{out} = 0.07972097658196313$   
Transformed:  $E_{out} = 0.07922272047832586$

The transformed improved the 5-versus-all out-of-sample performance by 0.6% Thus, e is true.

The answer is [e]

problem 10

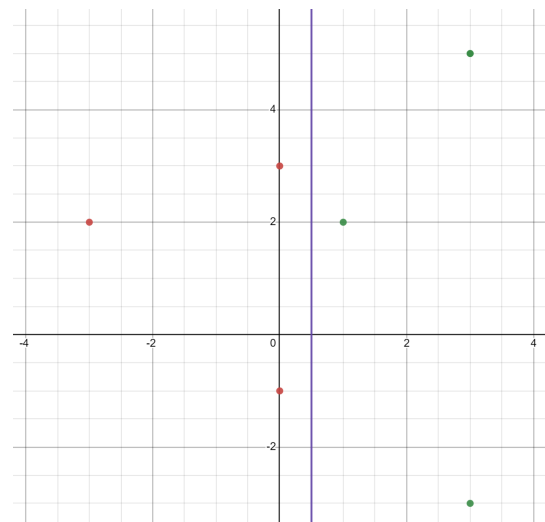
After running the code attached, the following values for  $E_{in}$  and  $E_{out}$  was produced:

Lambda: 0.01  
 $E_{out}$  : 0.0283  
 $E_{in}$  : 0.0045  
Lambda: 1  
 $E_{out}$  : 0.0259  
 $E_{in}$  : 0.0051

It follows that a is correct. As you can see, when  $\lambda = 1$ , the  $E_{in}$  was higher than when  $\lambda = 0.01$ , but the  $E_{out}$  was lower. This is a sign of over fitting. In this case, we would prefer the classifier for  $\lambda = 1$ , as this results in a lower  $E_{out}$ .

The answer is [a]

problem 11



When the transformed points are graphed, we see that each classifier lies on a different side of the line  $x = 0.5$ . Thus, only the  $z_1$  value for each point matters. To the right of the line, the points are classified as  $+1$ , and  $-1$  on the other side, thus,  $w_{z_1}$  is  $1$ , and  $w_{z_2}$  is  $0$ . The bias is  $-0.5$  as this is the shift of the separating plane. Thus,  $w_1 = 1, w_2 = 0, b = -0.5$ .

The answer is [c]

## problem 12

See code attached. I get 5 support vectors.

The answer is [c]

## problem 13

See code attached. For 100 iterations, the data set generated was never not separable by the RBF kernel. (0% of the time)

The answer is [a]

## problem 14

See code attached. For  $K = 9$ , and  $\gamma = 1.5$ , the kernel form had a lower  $E_{out}$  than the regular form 78% of the time.

The answer is [e]

## problem 15

See code attached. For  $K = 12$ , and  $\gamma = 1.5$ , the kernel form had a lower  $E_{out}$  than the regular form 62.6% of the time.

The answer is [d]

## problem 16

See code attached. The frequency of each happening (over 1000 iterations) were as follows:

- (a) 190
- (b) 85
- (c) 78
- (d) 420
- (e) 16

The answer is [d]

## problem 17

See code attached. The frequency of each happening (over 1000 iterations) were as follows:

- (a) 138
- (b) 174
- (c) 284
- (d) 173
- (e) 12

The answer is [c]

## problem 18

See attached code. When run over 1000 iterations using  $K = 0$ , and  $\gamma = 1.5$ , the regular RBF achieved  $E_{in} = 0$ , 1.4% of the time.

The answer is [a]

## problem 19

Since  $f$  is unknown across the interval  $[0, 1]$ , and since our random selection had a heart attack, we can describe the posterior probability. Lets say  $h$  is closer to 0,  $\mathcal{P}(h = f)$  will be much smaller than when  $h$  is near 1. Thus the posterior probability is not uniform. However, as each successive step (regardless of step size) towards 1 from 0 will result in  $\mathcal{P}(h = f)$  to also increase by the same step size. Thus, the posterior probability ( $\mathcal{P}(h = f)$ ) is linear over  $f \in [0, 1]$

The answer is [b]

## problem 20

Since we are using mean-squared error, consider some point  $\mathbf{x}$  that produces  $E_{g_1}$  and  $E_{g_2}$ . Since  $g$  is the aggregate of  $g_1$  and  $g_2$ , we know that  $g(\mathbf{x})$  will lie somewhere between  $g_1(\mathbf{x})$  and  $g_2(\mathbf{x})$ . Thus,  $E_g(\mathbf{x})$  will also lie between  $E_{g_1}$  and  $E_{g_2}$ . Thus,  $E_{out}(g)$  cannot be worse than the average of  $E_{out}(g_1)$  and  $E_{out}(g_2)$ .

The answer is [c]

The following code was used to answer questions 7-9:

```
import numpy

def versus_all(n, data_set):
    #splits data and classification info
    data = [[1, data_set[x][1], data_set[x][2], data_set[x][1]*data_set[x][2], data_set[x][1]*data_set[x][2]]
    classifier = [1 if x[0] == n else -1 for x in data_set]
    return classifier, data

lam = 1

for i in [0, 1, 2, 3, 4, 5, 6, 7, 8, 9]:
    class_i, data_i = versus_all(i, train_data)

    X_points = numpy.array(data_i)
    class_array = numpy.array(class_i)
    first_half = numpy.linalg.pinv(numpy.add(numpy.matmul(X_points.T, X_points), lam*numpy.identity(6)))
    w = numpy.matmul(first_half, numpy.matmul(X_points.T, class_array))

    true_class, test_data_i = versus_all(i, test_data)

    test_class_calc = [sign(numpy.dot(test_data_i[x], w)) for x in range(0, len(test_data_i))]

    E_out = (count_diff(true_class, test_class_calc))/len(true_class)
```

The following code was used to answer question 10:

```
import numpy

lam = 1

train_1v5 = []
test_1v5 = []

for x in train_data:
    if (x[0] == 1 or x[0] == 5):
        train_1v5.append(x)

for x in test_data:
    if (x[0] == 1 or x[0] == 5):
        test_1v5.append(x)

for lam in [0.01, 1]:
    class_i, data_i = versus_all(1, train_1v5)

    X_points = numpy.array(data_i)
    class_array = numpy.array(class_i)
    first_half = numpy.linalg.pinv(numpy.add(numpy.matmul(X_points.T, X_points), lam*numpy.identity(6)))
    w = numpy.matmul(first_half, numpy.matmul(X_points.T, class_array))

    true_class, test_data_i = versus_all(1, test_1v5)

    test_class_calc = [sign(numpy.dot(test_data_i[x], w)) for x in range(0, len(test_data_i))]

    E_out = (count_diff(true_class, test_class_calc))/len(true_class)
    E_in = count_diff([sign(numpy.dot(data_i[x], w)) for x in range(0, len(test_data_i))], class_i)/len

    print("Lamda: "+str(lam))
    print("E_out: "+str(E_out))
    print("E_in: "+str(E_in))
```

The following code was used to answer question 12:

```
from sklearn import svm
import numpy
from math import inf

data = [[1, 0], [0, 1], [0, -1], [-1, 0], [0, 2], [0, -2], [-2, 0]]
y = [-1, -1, -1, 1, 1, 1, 1]

clf = svm.SVC(kernel="poly", C=inf, gamma=1, coef0=1)
clf.fit(data, y)
print(clf.support_vectors_)
```

The following code was used to answer questions 13-15:

```
import numpy
from random import uniform
from sklearn import svm, cluster
import math

N = 100
gam = 1.5
clusters = 12
iterations = 1000

kernel_wins = 0

for i in range(0, iterations):
    data = [[uniform(-1, 1), uniform(-1, 1)] for x in range(0, N)]
    data_out = [[uniform(-1, 1), uniform(-1, 1)] for x in range(0, N)]

    def f(point):
        return sign(point[1] - point[0] + 0.25 * numpy.sin(numpy.pi*point[0]))

    class_data = [f(point) for point in data]
    class_data_out = [f(point) for point in data_out]

    clf = svm.SVC(kernel="rbf", C=math.inf, gamma=gam, coef0 = 1)
    clf.fit(data, class_data)

    Kmeans = cluster.KMeans(n_clusters=clusters, init='random').fit(data)

    for i in range(0, clusters):
        if (i not in Kmeans.labels_):
            iterations += 1
            continue

    phi = numpy.array(build_phi(data, clusters, Kmeans.cluster_centers_))
    class_data = numpy.array(class_data)

    w = numpy.matmul(numpy.matmul(numpy.linalg.pinv(numpy.matmul(phi.T, phi)), phi.T), class_data.T)

    #KMEANS OUT PREDICTION
    phi_out = numpy.array(build_phi(data_out, clusters, Kmeans.cluster_centers_))

    y_out = numpy.matmul(phi_out, w.T)
    y_out = [sign(x) for x in y_out]

    E_out_Kmeans = count_diff(y_out, class_data_out)/len(y_out)
    E_out_RBF = count_diff(clf.predict(data_out), class_data_out)/len(data_out)

    if E_out_RBF < E_out_Kmeans:
        kernel_wins += 1/iterations

def build_phi(points, clusters, centers):
    #BUILD PHI TO GET WEIGHTS
    phi = []
    for m in range(0, len(points)):
        point = points[m]
        row = []
        for n in range(0, clusters):
            center = centers[n]
            temp = numpy.exp(gam*-1 * numpy.linalg.norm(numpy.subtract(point, center))* numpy.linalg.noise)
            row.append(temp)
        phi.append(row)
    return phi

print(kernel_wins)
```



The following code was used to answer questions 16-18:

```
import numpy
from random import uniform
from sklearn import svm, cluster
import math

N = 100
gam = 1.5
#clusters = 9
iterations = 1000

answer = [0, 0, 0, 0, 0]

E_i = 0
for i in range(0, iterations):
    compare = []
    for clusters in [9]:
        data = [[uniform(-1, 1), uniform(-1, 1)] for x in range(0, N)]
        data_out = [[uniform(-1, 1), uniform(-1, 1)] for x in range(0, N)]

        def f(point):
            return sign(point[1] - point[0] + 0.25 * numpy.sin(numpy.pi*point[0]))

        class_data = [f(point) for point in data]
        class_data_out = [f(point) for point in data_out]

        """clf = svm.SVC(kernel="rbf", C=math.inf, gamma=gam, coef0 = 1)
        clf.fit(data, class_data) """

        Kmeans = cluster.KMeans(n_clusters=clusters, init='random').fit(data)

        for i in range(0,clusters):
            if (i not in Kmeans.labels_):
                iterations += 1
                continue

        phi = numpy.array(build_phi(data, clusters, Kmeans.cluster_centers_))
        class_data = numpy.array(class_data)

        w = numpy.matmul(numpy.matmul(numpy.linalg.pinv(numpy.matmul(phi.T, phi)), phi.T), class_data.T)

        #KMEANS OUT PREDICTION
        phi_out = numpy.array(build_phi(data_out, clusters, Kmeans.cluster_centers_))

        y_out = numpy.matmul(phi_out, w.T)
        y_out = [sign(x) for x in y_out]

        y_in = numpy.matmul(phi, w.T)
        y_in = [sign(x) for x in y_in]

        E_out_Kmeans = count_diff(y_out, class_data_out)/len(y_out)
        E_in_Kmeans = count_diff(y_in, class_data)/len(y_in)
        compare.append(E_out_Kmeans)
        compare.append(E_in_Kmeans)

    if compare[1] == 0:
        E_i += 1/iterations

    E_in_goes_up = compare[1] < compare[3]
    E_in_goes_down = compare[1] > compare[3]
    E_out_goes_up = compare[0] < compare[2]
    E_out_goes_down = compare[0] > compare[2]
    no_change = bool((compare[1] == compare[3]) & (compare[0] == compare[2]))
    if E_in_goes_down & E_out_goes_up:
        answer[0] += 1

    if E_in_goes_up & E_out_goes_down:
        answer[1] += 1

    if E_in_goes_up & E_out_goes_up:
        answer[2] += 1
```

```
if E_in_goes_down & E_out_goes_down:
    answer[3] += 1

if no_change:
    answer[4] += 1
print(answer)
```